

Classwork 2: Variable Selection and Model Optimization

Introduction to Predictive Analytics in R

Prof. Asc. Endri Raco, Ph.D.

Polytechnic University of Tirana

November 2025

Section 1

Classwork Overview

Today's Challenge: Optimize Your Model

Select the Best Variables for Prediction

Scenario: You're back at the bank! They loved your first model, but now they want you to:

- ① Improve model performance using variable selection
- ② Avoid overfitting
- ③ Identify the optimal number of predictors

Time: 30 minutes

Work: Individually or in pairs

Learning Objectives

By the end of this classwork, you will:

- ① Implement forward stepwise variable selection
- ② Split data into training and test sets
- ③ Calculate and interpret AUC
- ④ Detect overfitting
- ⑤ Choose the optimal number of variables
- ⑥ Visualize model performance

Dataset: Enhanced Bank Marketing

Target Variable: subscribed (yes = 1, no = 0)

Available Predictors (12 variables):

- age: Client's age
- balance: Average yearly balance (euros)
- duration: Last contact duration (seconds)
- campaign: Number of contacts this campaign
- previous: Number of previous contacts
- pdays: Days since last contact
- job_*: Job type indicators (3 variables)
- education_*: Education level indicators (3 variables)

Getting Started

Step 1: Setup Your Environment

Create a new R script called `classwork2.R`

```
# Load required libraries
library(tidyverse)
library(pROC)
library(caret)

# Set seed for reproducibility
set.seed(456)

# Load the data
bank_data <- read_csv("bank_marketing_v2.csv")
```


Section 2

Part 1: Data Preparation (5 minutes)

Task 1.1: Explore the Data

```
# View structure  
glimpse(bank_data)  
  
# Check dimensions  
cat("Rows:", nrow(bank_data), "\n")  
cat("Columns:", ncol(bank_data), "\n")  
  
# View first few rows  
head(bank_data)
```

Question 1: How many candidate predictors do you have?

Task 1.2: Check Target Distribution

```
# Target variable summary
table(bank_data$subscribed)

# Subscription rate
sub_rate <- mean(bank_data$subscribed)
cat("Subscription rate:",
    round(sub_rate * 100, 2), "%\n")

# Visualize
ggplot(bank_data, aes(x = factor(subscribed))) +
  geom_bar(fill = "steelblue") +
  labs(title = "Target Distribution",
       x = "Subscribed", y = "Count") +
  theme_minimal()
```

Task 1.3: Train/Test Split

```
# Create 60/40 split
train_index <- createDataPartition(
  bank_data$subscribed,
  p = 0.6,
  list = FALSE
)

# Split data
train_data <- bank_data[train_index, ]
test_data <- bank_data[-train_index, ]
```

Task 1.4: Verify Split

```
# Check sizes
cat("Training size:", nrow(train_data), "\n")
cat("Test size:", nrow(test_data), "\n")

# Check balance
cat("Train subscription rate:",
    round(mean(train_data$subscribed)*100, 2),
    "%\n")
cat("Test subscription rate:",
    round(mean(test_data$subscribed)*100, 2),
    "%\n")
```

Question 2: Are the training and test sets balanced?

Section 3

Part 2: Create Helper Functions (8 minutes)

Task 2.1: AUC Calculation Function

```
# Function to calculate AUC
calculate_auc <- function(variables,
                           target,
                           data) {

  # Build formula
  formula_str <- paste(target, "~",
                        paste(variables,
                               collapse = " + "))
  formula <- as.formula(formula_str)

  # Continue on next slide...
}
```

Task 2.1: AUC Function (continued)

```
calculate_auc <- function(variables,
                           target,
                           data) {

  formula_str <- paste(target, "~",
                        paste(variables,
                               collapse = " + "))
  formula <- as.formula(formula_str)

  # Fit model
  model <- glm(formula,
                data = data,
                family = binomial)

  # Get predictions
  preds <- predict(model, type = "response")

  # Calculate AUC
  auc_val <- auc(roc(data[[target]]), preds)}
```

Task 2.2: Test AUC Function

```
# Test with a single variable
test_auc <- calculate_auc(
  variables = "age",
  target = "subscribed",
  data = train_data
)
cat("AUC with age only:",
    round(test_auc, 3), "\n")
```

Question 3: What is the AUC using only age?

Task 2.3: Find Next Best Variable Function

```
# Function to find next best variable
find_next_best <- function(current_vars,
                             candidate_vars,
                             target,
                             data) {

  best_auc <- -1
  best_var <- NULL

  for (var in candidate_vars) {
    # Test adding this variable
    test_vars <- c(current_vars, var)

    # Continue on next slide...
  }
}
```

Task 2.3: Find Next Best (continued)

```
find_next_best <- function(current_vars,
                             candidate_vars,
                             target, data) {
  best_auc <- -1
  best_var <- NULL

  for (var in candidate_vars) {
    test_vars <- c(current_vars, var)
    test_auc <- calculate_auc(test_vars,
                               target, data)

    # Update if better
    if (test_auc > best_auc) {
      best_auc <- test_auc
      best_var <- var
    }
  }

  return(list(variable = best_var,
             auc = best_auc))
```

Task 2.4: Test Find Next Best

```
# Define candidate variables
candidates <- c("age", "balance", "duration",
               "campaign", "previous")

# Find best first variable
result <- find_next_best(
  current_vars = c(),
  candidate_vars = candidates,
  target = "subscribed",
  data = train_data
)
cat("Best variable:", result$variable, "\n")
cat("AUC:", round(result$auc, 3), "\n")
```

Question 4: Which variable gives the highest AUC?

Section 4

Part 3: Forward Stepwise Selection (10 minutes)

Task 3.1: Setup Variables

```
# All candidate predictors
all_candidates <- c(
  "age", "balance", "duration",
  "campaign", "previous", "pdays",
  "job_management", "job_technician",
  "job_blue-collar",
  "education_primary", "education_secondary",
  "education_tertiary"
)

# Initialize
current_vars <- c()
candidate_vars <- all_candidates
max_vars <- 8
```

Task 3.2: Forward Selection Loop

```
# Storage for results
selected_vars <- c()
train_aucs <- c()
test_aucs <- c()

# Run forward selection
for (i in 1:max_vars) {

  # Find next best on training data
  result <- find_next_best(
    current_vars = current_vars,
    candidate_vars = candidate_vars,
    target = "subscribed",
    data = train_data
  )

  # Continue on next slide...
}
```

Task 3.2: Forward Selection (continued)

```
for (i in 1:max_vars) {  
  
  result <- find_next_best(  
    current_vars, candidate_vars,  
    "subscribed", train_data  
  )  
  
  # Update variables  
  next_var <- result$variable  
  current_vars <- c(current_vars, next_var)  
  candidate_vars <- setdiff(candidate_vars,  
    next_var)  
  
  # Continue on next slide...  
}
```

Task 3.2: Forward Selection (final part)

```
for (i in 1:max_vars) {  
  # ... previous code ...  
  
  current_vars <- c(current_vars, next_var)  
  candidate_vars <- setdiff(candidate_vars,  
                            next_var)  
  
  # Calculate AUCs  
  train_auc <- calculate_auc(current_vars,  
                                "subscribed",  
                                train_data)  
  test_auc <- calculate_auc(current_vars,  
                             "subscribed",  
                             test_data)  
  
  # Store results  
  selected_vars <- c(selected_vars, next_var)  
  train_aucs <- c(train_aucs, train_auc)  
  test_aucs <- c(test_aucs, test_auc)
```

Task 3.3: View Results

```
# Create results dataframe
results <- data.frame(
  step = 1:max_vars,
  variable = selected_vars,
  train_auc = train_aucs,
  test_auc = test_aucs,
  gap = train_aucs - test_aucs
)
# Display results
print(results)
```

Question 5: At which step does test AUC peak?

Section 5

Part 4: Visualization & Analysis (7 minutes)

Task 4.1: Plot Train vs Test AUC

```
# Reshape for plotting
results_long <- results %>%
  pivot_longer(cols = c(train_auc, test_auc),
               names_to = "dataset",
               values_to = "auc")

# Create plot
ggplot(results_long, aes(x = step,
                           y = auc,
                           color = dataset)) +
  geom_line(size = 1.2) +
  geom_point(size = 3)
```

Task 4.1: Enhance the Plot

```
ggplot(results_long, aes(x = step,
                          y = auc,
                          color = dataset)) +
  geom_line(size = 1.2) +
  geom_point(size = 3) +
  scale_color_manual(
    values = c("train_auc" = "red",
              "test_auc" = "blue"),
    labels = c("Training", "Test"))
) +
  labs(title = "Model Performance: Train vs Test",
       x = "Number of Variables",
       y = "AUC",
       color = "Dataset") +
  theme_minimal() +
  theme(legend.position = "top")
```

Task 4.2: Identify Optimal Model

```
# Find step with highest test AUC
optimal_step <- which.max(test_aucs)

cat("Optimal number of variables:",
    optimal_step, "\n")
cat("Optimal test AUC:",
    round(test_aucs[optimal_step], 3), "\n")

# Get optimal variables
optimal_vars <- selected_vars[1:optimal_step]
cat("\nOptimal variables:\n")
print(optimal_vars)
```

Question 6: What are your optimal variables?

Task 4.3: Analyze Overfitting

```
# Calculate gap (overfitting indicator)
results$gap <- results$train_auc -
            results$test_auc

# Plot the gap
ggplot(results, aes(x = step, y = gap)) +
  geom_line(color = "orange", size = 1.2) +
  geom_point(size = 3, color = "orange") +
  geom_hline(yintercept = 0.05,
             linetype = "dashed",
             color = "red") +
  labs(title = "Overfitting Gap",
       x = "Number of Variables",
       y = "Train AUC - Test AUC") +
  theme_minimal()
```

Task 4.4: Build Final Model

```
# Build model with optimal variables
final_formula <- as.formula(
  paste("subscribed ~",
        paste(optimal_vars, collapse = " + ")))
)

final_model <- glm(final_formula,
                     data = train_data,
                     family = binomial)

# View coefficients
summary(final_model)
```

Task 4.5: Final Evaluation

```
# Predict on test set
test_preds <- predict(final_model,
                      newdata = test_data,
                      type = "response")

# Calculate final metrics
final_roc <- roc(test_data$subscribed,
                   test_preds)
final_auc <- auc(final_roc)

cat("Final Test AUC:",
    round(final_auc, 3), "\n")

# Plot ROC curve
plot(final_roc, main = "ROC Curve - Final Model",
      col = "blue", lwd = 2)
```

Task 4.6: Compare with Full Model

```
# Build model with ALL variables
full_model <- glm(
  subscribed ~ .,
  data = train_data,
  family = binomial
)

# Evaluate on test set
full_preds <- predict(full_model,
                      newdata = test_data,
                      type = "response")

full_auc <- auc(roc(test_data$subscribed,
                     full_preds))

cat("Full model AUC:", round(full_auc, 3), "\n")
cat("Optimal model AUC:",
    round(final_auc, 3), "\n")
```


Section 6

Bonus Challenges

Bonus Challenge 1: Backward Selection

```
# Start with all variables
current_vars_back <- all_candidates
removed_vars <- c()
backward_aucs <- c()

# Iteratively remove worst variable
while(length(current_vars_back) > 3) {

  worst_auc <- Inf
  worst_var <- NULL

  # Try removing each variable
  for (var in current_vars_back) {
    test_vars <- setdiff(current_vars_back,
                          var)
    test_auc <- calculate_auc(test_vars,
                               "subscribed",
                               train_data)

  # Keep track of worst (causes smallest drop)
  # To be continued in the next slide
```

Bonus Challenge 2: Cross-Validation

```
# 5-fold cross-validation
library(caret)

# Define training control
train_control <- trainControl(
  method = "cv",
  number = 5,
  summaryFunction = twoClassSummary,
  classProbs = TRUE
)

# Train with cross-validation
cv_model <- train(
  as.factor(subscribed) ~
    duration + balance + previous,
  data = train_data,
  method = "glm",
  family = "binomial",
  trControl = train_control,
  metric = "ROC"
)
```

Bonus Challenge 3: Feature Importance

```
# Calculate variable importance
# Using coefficient magnitude

coefs <- coef(final_model)[-1] # Remove intercept
importance <- abs(coefs)

# Create dataframe
var_importance <- data.frame(
  variable = names(importance),
  importance = importance
) %>%
  arrange(desc(importance))

# Visualize
ggplot(var_importance,
       aes(x = reorder(variable, importance),
            y = importance)) +
  geom_col(fill = "steelblue") +
  coord_flip() +
  labs(title = "Variable Importance",
       subtitle = "Using coefficient magnitude")
```


Section 7

Summary & Deliverables

What You Accomplished

- ① ✓ Implemented forward stepwise selection
- ② ✓ Split data properly for validation
- ③ ✓ Calculated AUC for multiple models
- ④ ✓ Detected and analyzed overfitting
- ⑤ ✓ Identified optimal number of variables
- ⑥ ✓ Built and evaluated final model

Submission Requirements

What to Submit

1. Complete R script (`classwork2.R`)
2. Brief report answering all 7 questions
3. Three visualizations:
 - Train vs Test AUC plot
 - Overfitting gap plot
 - ROC curve for final model
4. Optional: Bonus challenge solutions

Deadline: Before next lecture

Format: PDF or Word document + R script

Answer Key Template

Question 1: Number of predictors = _____

Question 2: Train/test balanced? _____

Question 3: AUC with age = _____

Question 4: Best first variable = _____

Question 5: Test AUC peaks at step = _____

Question 6: Optimal variables: _____

Question 7: Better model: _____

Key Insights to Report

In your report, discuss:

- ① **Variable Selection Results** Which variables were selected? Why do they make sense?
- ② **Overfitting Analysis** At what point did overfitting occur? How did you detect it?
- ③ **Model Comparison** How does your optimal model compare to using all variables?
- ④ **Business Recommendations** Which variables should the bank focus on collecting?

Grading Rubric

Component	Points
Code runs without errors	25
All 7 questions answered correctly	35
Three required visualizations	20
Code quality and comments	10
Written analysis and insights	10
Bonus challenges	+5 each
Total	100

Common Issues & Solutions

Issue 1: Functions not working

Solution: Check that pROC and caret are loaded

Issue 2: AUC values seem too low

Solution: Check target variable is numeric (0/1)

Issue 3: Loop taking too long

Solution: Reduce max_vars or use smaller dataset

Tips for Success

Efficiency Tips

- Test functions with small examples first - Use `head()` to preview data before full analysis - Comment your code as you go - Save intermediate results

Analysis Tips

- Look for the "elbow" in test AUC curve - Small gaps (< 0.05) between train/test are good - Business relevance matters too!

Getting the Data

Option 1: Download from course website

- URL: www.pythontesting.net/data/bank_marketing_v2.csv

Option 2: Use the dataset from Classwork 1

Option 3: Generate sample data

```
set.seed(456)
n <- 2000
bank_data <- data.frame(
  age = rnorm(n, 40, 12),
  balance = rnorm(n, 1500, 3000),
  duration = rnorm(n, 250, 150),
  campaign = rpois(n, 2),
  previous = rpois(n, 0.5),
  pdays = rpois(n, 40),
  subscribed = rbinom(n, 1, 0.15)
)
# Add job and education dummies...
```

Help Resources

During Classwork:

- Raise your hand for assistance
- Collaborate with neighbors (but write your own code!)
- Refer to lecture slides

Reference Materials:

- `?glm` - Logistic regression help
- `?roc` - ROC curve documentation
- `?createDataPartition` - Data splitting

Office Hours: Tuesday & Thursday, 2-4 PM

Learning Outcomes Check

After completing this classwork, you should:

- Understand why variable selection matters
- Implement forward stepwise selection from scratch
- Properly validate models using train/test split
- Recognize signs of overfitting
- Choose optimal model complexity
- Communicate results to stakeholders

Next Steps

In Next Lecture:

- Cross-validation techniques
- Regularization (Lasso, Ridge)
- Advanced feature engineering
- Handling imbalanced data

For Next Time:

- Complete this classwork
- Review AUC interpretation
- Read Chapter 3: Model Evaluation
- Bring questions!

Reflection Questions

Think about (not graded):

- ① Why is forward selection “greedy”? What are its limitations?
- ② Could you have gotten different results with backward selection?
- ③ How would you explain your variable selection to a non-technical manager?
- ④ What if collecting data for some variables is very expensive?

Quick Reference: Key Functions

```
# Data splitting
createDataPartition(y, p = 0.6, list = FALSE)

# AUC calculation
roc_obj <- roc(actual, predicted)
auc(roc_obj)

# Model building
glm(formula, data, family = binomial)
predict(model, newdata, type = "response")

# Visualization
ggplot(data, aes(x, y, color)) +
  geom_line() + geom_point()
```

Advanced: Comparing Multiple Methods

```
# Compare forward, backward, and all variables
methods <- data.frame(
  method = c("Forward", "Backward", "All"),
  n_vars = c(length(optimal_vars),
             length(backward_vars),
             length(all_candidates)),
  test_auc = c(forward_auc,
               backward_auc,
               full_auc)
)
ggplot(methods, aes(x = n_vars,
                     y = test_auc,
                     color = method)) +
  geom_point(size = 4) +
  geom_line() +
  labs(title = "Method Comparison") +
  theme_minimal()
```

Ready to Start!

Let's Begin!

You have 30 minutes

*Focus on understanding the concepts,
not just getting the right answer*

Good luck!

Need Help?

Remember:

Neighbors can be great resources

Lecture slides have all the code you need

Don't spend more than 5 minutes stuck on one problem!