

# Classwork 3: Data Preprocessing

## Categorical Encoding, Missing Values, and Transformations

Prof. Asc. Endri Raco, Ph.D.

November 2025

## Overview

**Objective:** Master data preprocessing techniques essential for building robust predictive models, including handling categorical variables, missing values, outliers, and variable transformations.

**Dataset:** `donor_transactions.csv` (same as Classwork 2)

**Submission:** R Markdown file (.Rmd) + knitted PDF

**Due Date:** November 27, 2025, 23:59

**Weight:** 15% of final grade

---

## Part 1: Creating Dummy Variables (20 points)

### Task 1.1: Single Categorical Variable (8 points)

You have a `segment` variable with three categories: Gold, Silver, Bronze.

```
# Load required libraries
library(tidyverse)
library(fastDummies)

# Load your basetable from Classwork 2
basetable <- read_csv("basetable_classwork2.csv")

# Examine the segment variable
table(basetable$segment)
```

**Tasks:**

- Create dummy variables for `segment` using the `fastDummies` package
- Ensure you drop one category to avoid multicollinearity
- Remove the original `segment` column
- Show the first 10 rows with `donor_id` and the new dummy variables

**Questions to answer:**

- Which category did you choose as reference? Why?
- How many dummy variables were created?
- What does it mean when all dummy variables equal 0?

**Deliverable:** Code + output table + written explanations

---

## Task 1.2: Multiple Categorical Variables (7 points)

Create dummy variables for multiple categorical columns simultaneously.

```
# Identify all categorical variables in your basetable
# Likely candidates: segment, country, gender, etc.

# Create dummies for all categorical variables at once
categorical_vars <- c("segment", "country") # Add more if you have them

basetable_with_dummies <- dummy_cols(
  basetable,
  select_columns = categorical_vars,
  remove_first_dummy = TRUE,
  remove_selected_columns = TRUE
)
```

### Tasks:

- Create dummy variables for at least 2 categorical variables
- Document which reference category was chosen for each variable
- Count total number of dummy columns created
- Verify no multicollinearity by checking correlation between dummies

### Deliverable:

- Summary table showing original categories and resulting dummies
- Explanation of reference category selection
- Verification that no perfect correlation exists

---

## Task 1.3: Manual Dummy Creation (5 points)

Create dummy variables manually to understand the process.

```
# Manual approach for segment variable
basetable$segment_Gold <- as.integer(basetable$segment == "Gold")
basetable$segment_Silver <- as.integer(basetable$segment == "Silver")
# Bronze is reference (both equal 0)

# Verify your manual dummies match the package output
```

### Questions:

- Why use `as.integer()` instead of just TRUE/FALSE?
- How would you manually create dummies for a variable with 5 categories?
- What happens if a donor's segment is NA?

### Deliverable:

Working code + answers to conceptual questions

---

## Part 2: Handling Missing Values (25 points)

### Task 2.1: Missing Value Assessment (7 points)

First, understand the extent and patterns of missing data.

```
# Count missing values per column
missing_summary <- data.frame(
  variable = names(basetable),
  n_missing = sapply(basetable, function(x) sum(is.na(x))),
  pct_missing = sapply(basetable, function(x)
    round(100 * sum(is.na(x)) / length(x), 2))
)

# Sort by percentage missing
missing_summary <- missing_summary[order(-missing_summary$pct_missing), ]
```

#### Tasks:

- Create the missing value summary table
- Identify variables with >10% missing values
- Create a visualization (bar chart) of missing percentages
- Discuss patterns: Are missing values random or systematic?

#### Deliverable:

- Missing value summary table
- Visualization of missing data
- Written analysis of missingness patterns

---

### Task 2.2: Imputation Strategy Selection (8 points)

Develop appropriate imputation strategies for different variable types.

```
# Strategy 1: Mean imputation for roughly normal variables
# Example: age
if (sum(is.na(basetable$age)) > 0) {
  mean_age <- mean(basetable$age, na.rm = TRUE)
  basetable$age_imputed <- basetable$age
  basetable$age_imputed[is.na(basetable$age_imputed)] <- mean_age
}

# Strategy 2: Median imputation for skewed variables
# Example: max_donation

# Strategy 3: Zero imputation for count variables
# Example: donations_last_year
```

#### Tasks:

- For EACH variable with missing values, justify your imputation strategy:
  - Mean: Use for approximately normal distributions
  - Median: Use for skewed distributions with outliers
  - Zero: Use when missing means “absence” (e.g., no donations)
  - Mode: Use for categorical variables
- Implement your chosen strategies for at least 5 variables

- c) Create “before and after” histograms for 3 numeric variables

**Deliverable:**

- Imputation strategy table with justifications
  - Code implementing all strategies
  - Before/after visualizations
  - Discussion of how imputation affects distributions
- 

### Task 2.3: Missing Value Indicators (5 points)

Create indicator variables that flag when values were missing.

```
# Create missing indicators BEFORE imputing
basetable$missing_email <- as.integer(is.na(basetable$email))
basetable$missing_phone <- as.integer(is.na(basetable$phone))
basetable$missing_age <- as.integer(is.na(basetable$age))

# Analyze: Are donors with missing email less likely to donate?
aggregate(target ~ missing_email, data = basetable, mean)
```

**Tasks:**

- a) Create missing indicators for at least 3 variables
- b) Calculate the donation rate for each missing indicator
- c) Identify which missing indicators are most predictive
- d) Explain why “missingness” itself might be informative

**Deliverable:**

- Code creating indicators
  - Donation rate comparison table
  - Top 3 most predictive missing indicators
  - Business interpretation of findings
- 

### Task 2.4: Advanced Imputation (5 points)

Implement a more sophisticated imputation method.

```
# Option 1: Group-based imputation
# Impute age based on segment
basetable <- basetable %>%
  group_by(segment) %>%
  mutate(age_imputed_group = ifelse(
    is.na(age),
    median(age, na.rm = TRUE),
    age
  )) %>%
  ungroup()

# Option 2: Predictive imputation using simple regression
# Use other variables to predict missing values
```

**Tasks:**

- a) Implement group-based imputation for at least 1 variable

- b) Compare simple median imputation vs. group-based imputation
- c) Calculate improvement in variance preservation

**Deliverable:**

- Code for advanced imputation
  - Comparison table of methods
  - Discussion of when to use each approach
- 

## Part 3: Handling Outliers (25 points)

### Task 3.1: Outlier Detection (8 points)

Identify outliers using multiple methods.

```
# Method 1: Visual inspection
par(mfrow = c(2, 2))
boxplot(basetable$mean_donation, main = "Mean Donation")
boxplot(basetable$max_donation, main = "Max Donation")
boxplot(basetable$age, main = "Age")
boxplot(basetable$days_since_last, main = "Days Since Last")

# Method 2: Statistical thresholds
# Define outliers as values beyond 3 SD
detect_outliers_sd <- function(x, n_sd = 3) {
  mean_x <- mean(x, na.rm = TRUE)
  sd_x <- sd(x, na.rm = TRUE)
  lower <- mean_x - n_sd * sd_x
  upper <- mean_x + n_sd * sd_x

  outliers <- x < lower | x > upper
  return(outliers)
}

# Method 3: IQR method
detect_outliers_iqr <- function(x) {
  q1 <- quantile(x, 0.25, na.rm = TRUE)
  q3 <- quantile(x, 0.75, na.rm = TRUE)
  iqr <- q3 - q1

  lower <- q1 - 1.5 * iqr
  upper <- q3 + 1.5 * iqr

  outliers <- x < lower | x > upper
  return(outliers)
}
```

**Tasks:**

- a) Create boxplots for at least 5 numeric variables
- b) Apply both SD method and IQR method to detect outliers
- c) Create a summary table comparing methods
- d) Identify which variables have the most outliers

**Deliverable:**

- Boxplot visualizations
  - Outlier detection summary table
  - Comparison of SD vs. IQR methods
  - Discussion of which method is more appropriate for your data
- 

### Task 3.2: Winsorization (8 points)

Apply winsorization to cap extreme values.

```
# Install DescTools if needed
# install.packages("DescTools")
library(DescTools)

# Winsorize at 5th and 95th percentiles
basetable$mean_donation_wins <- Winsorize(
  basetable$mean_donation,
  probs = c(0.05, 0.95),
  na.rm = TRUE
)

# Compare distributions
par(mfrow = c(1, 2))
hist(basetable$mean_donation, main = "Original",
      xlab = "Mean Donation", breaks = 30)
hist(basetable$mean_donation_wins, main = "Winsorized",
      xlab = "Mean Donation", breaks = 30)
```

Tasks:

- a) Winsorize at least 3 numeric variables
- b) Use both 5%/95% and 1%/99% winsorization levels
- c) Create before/after histograms for each variable
- d) Calculate summary statistics (mean, SD, min, max) before and after

Deliverable:

- Winsorization code for multiple variables
  - Before/after visualizations
  - Summary statistics comparison table
  - Recommendation on which percentile levels to use
- 

### Task 3.3: Standard Deviation Method (9 points)

Implement the mean  $\pm 3$  SD capping method.

```
# Function to cap outliers at mean ± n SD
cap_outliers_sd <- function(x, n_sd = 3) {
  mean_x <- mean(x, na.rm = TRUE)
  sd_x <- sd(x, na.rm = TRUE)

  lower_limit <- mean_x - n_sd * sd_x
  upper_limit <- mean_x + n_sd * sd_x

  # Cap values
```

```

x_capped <- pmin(pmax(x, lower_limit), upper_limit)

return(x_capped)
}

# Apply to variables
basetable$mean_donation_capped <- cap_outliers_sd(
  basetable$mean_donation,
  n_sd = 3
)

```

**Tasks:**

- a) Implement the SD capping function
- b) Apply to at least 3 variables with outliers
- c) Experiment with 2 SD, 3 SD, and 4 SD thresholds
- d) Compare SD method vs. Winsorization on the same variable
- e) Analyze how many values were capped at each threshold

**Deliverable:**

- Capping function implementation
  - Comparison of different SD thresholds
  - Winsorization vs. SD method comparison
  - Recommendation table: which method for which variable
  - Discussion of trade-offs
- 

## Part 4: Variable Transformations (20 points)

### Task 4.1: Log Transformation (7 points)

Apply logarithmic transformation to right-skewed variables.

```

# Identify skewed variables
skewness_check <- sapply(
  basetable[, sapply(basetable, is.numeric)],
  function(x) {
    library(e1071)
    skewness(x, na.rm = TRUE)
  }
)

# Variables with skewness > 1 are candidates for log transform
print(sort(skewness_check, decreasing = TRUE))

# Apply log transformation
# Handle zeros with log1p (log(x + 1))
basetable$log_mean_donation <- log1p(basetable$mean_donation)
basetable$log_max_donation <- log1p(basetable$max_donation)

```

**Tasks:**

- a) Calculate skewness for all numeric variables
- b) Identify variables with absolute skewness > 1
- c) Apply log transformation to at least 3 skewed variables

- d) Create before/after histograms showing improvement
- e) Verify that transformed variables have lower skewness

**Deliverable:**

- Skewness summary table
  - Log transformation code
  - Before/after distribution comparisons (histograms + QQ plots)
  - New skewness values showing improvement
- 

### Task 4.2: Other Transformations (6 points)

Explore alternative transformation methods.

```
# Square root transformation (milder than log)
basetable$sqrt_mean_donation <- sqrt(basetable$mean_donation)

# Inverse transformation
basetable$inv_recency <- 1 / (basetable$days_since_last + 1)

# Box-Cox transformation (finds optimal power)
library(MASS)
bc <- boxcox(mean_donation ~ 1, data = basetable,
              lambda = seq(-2, 2, 0.1))
optimal_lambda <- bc$x[which.max(bc$y)]
```

**Tasks:**

- a) Apply square root transformation to 2 variables
- b) Create inverse transformation for recency (so recent = high value)
- c) Implement Box-Cox transformation for 1 variable
- d) Compare all transformation methods on the same variable

**Deliverable:**

- Code for all transformation types
  - Comparison table: original vs. log vs. sqrt vs. Box-Cox
  - Visualization showing distribution improvements
  - Recommendations for which transformation to use when
- 

### Task 4.3: Transformation Impact Analysis (7 points)

Analyze how transformations affect model performance.

```
# Simple logistic regression comparison
# Model 1: Original variables
model_original <- glm(
  target ~ mean_donation + days_since_last,
  data = basetable,
  family = "binomial"
)

# Model 2: Transformed variables
model_transformed <- glm(
  target ~ log_mean_donation + inv_recency,
```

```

    data = basetable,
    family = "binomial"
)

# Compare AIC
AIC(model_original)
AIC(model_transformed)

```

**Tasks:**

- a) Fit logistic regression with original variables
- b) Fit logistic regression with transformed variables
- c) Compare model fit statistics (AIC, deviance)
- d) Compare coefficient interpretability
- e) Calculate predicted probabilities for both models

**Deliverable:**

- Model comparison table (AIC, deviance, pseudo-R<sup>2</sup>)
  - Coefficient comparison and interpretation
  - Discussion: Is the improvement worth the interpretation complexity?
- 

## Part 5: Creating Interaction Features (10 points)

### Task 5.1: Meaningful Interactions (6 points)

Create interaction terms that make business sense.

```

# Frequency x Recency interaction
basetable$freq_recency <-
  basetable$donation_count / (basetable$days_since_last + 1)

# Frequency x Monetary interaction
basetable$freq_monetary <-
  basetable$donation_count * basetable$mean_donation

# Three-way RFM interaction
basetable$rfm_score <-
  (1 / (basetable$days_since_last + 1)) *
  basetable$donation_count *
  basetable$mean_donation

```

**Tasks:**

- a) Create at least 5 meaningful interaction features
- b) Explain the business logic behind each interaction
- c) Examine the distribution of interaction features
- d) Identify top 10 donors by RFM score

**Deliverable:**

- Interaction creation code
- Business justification for each interaction
- Distribution summary statistics
- Top/bottom 10 donors comparison

---

## Task 5.2: Interaction Impact (4 points)

Test whether interactions improve model performance.

```
# Model without interactions
model_no_int <- glm(
  target ~ donation_count + days_since_last + mean_donation,
  data = basetable,
  family = "binomial"
)

# Model with interactions
model_with_int <- glm(
  target ~ donation_count + days_since_last + mean_donation +
    freq_recency + rfm_score,
  data = basetable,
  family = "binomial"
)

# Compare models
anova(model_no_int, model_with_int, test = "Chisq")
```

### Tasks:

- Fit models with and without interactions
- Perform likelihood ratio test
- Compare predictive accuracy on holdout set
- Identify which interactions are most important

### Deliverable:

- Model comparison statistics
  - Likelihood ratio test results
  - Discussion of which interactions add value
- 

## Grading Rubric

Component	Points	Criteria
Part 1: Dummy Variables	20	Correct implementation, proper multicollinearity handling, clear documentation
Part 2: Missing Values	25	Appropriate strategy selection, thorough analysis, missing indicators
Part 3: Outliers	25	Multiple detection methods, proper capping techniques, method comparison
Part 4: Transformations	20	Appropriate transformations, impact analysis, clear visualizations
Part 5: Interactions	10	Business-motivated interactions, performance evaluation

Component	Points	Criteria
<b>Code Quality</b>	<b>10</b>	Clean, well-commented, reproducible
<b>Report Quality</b>	<b>10</b>	Clear explanations, professional presentation
<b>Total</b>	<b>120</b>	(100 + 20 quality points)

---

## Submission Guidelines

### What to Submit

1. **R Markdown file** (.Rmd) with all code, analysis, and explanations
2. **Knitted PDF** showing all output, visualizations, and tables
3. **Processed dataset** (CSV) with all new features created

### File Naming Convention

Classwork3\_FirstnameLastname.Rmd

Classwork3\_FirstnameLastname.pdf

basetable\_preprocessed\_FirstnameLastname.csv

### Submission Platform

Upload to course portal under “Assignments > Classwork 3”

---

## Tips for Success

1. **Document your decisions** - Explain WHY you chose each method
  2. **Compare methods** - Don't just apply one technique, compare alternatives
  3. **Visualize everything** - Before/after plots reveal transformation impact
  4. **Think about business** - Do your preprocessing choices make sense?
  5. **Check for errors** - Look for Inf, NaN, NA after transformations
  6. **Save intermediate results** - Keep original variables for comparison
  7. **Use functions** - Write reusable code for repetitive tasks
  8. **Test on subsets first** - Verify methods work before applying to all data
- 

## Frequently Asked Questions

**Q: Should I apply all transformations to all variables?**

A: No! Only transform variables that need it. Check distribution first.

**Q: What if my variable has no missing values?**

A: Great! Document that and move to next variable. No imputation needed.

**Q: How do I know if winsorization or SD method is better?**

A: Check your data distribution. Normal data → SD method. Skewed data → winsorization.

**Q: Can I create more than 5 interactions?**

A: Yes, but be selective. Focus on interactions with business justification.

**Q: What if log transformation creates Inf values?**

A: Use `log1p()` instead of `log()`, or add a small constant before logging.

**Q: Should I remove outliers or cap them?**

A: Generally cap them (winsorize). Only remove if they're data errors.

**Q: How many dummy variables should segment create?**

A: If segment has 3 categories, create 2 dummies (drop 1 for reference).

---

## Academic Integrity

- You may discuss concepts with classmates
- You must write your own code and explanations
- Copying code from others = automatic zero
- Using AI for debugging = OK
- Using AI to write entire solutions = NOT OK

**The goal is learning, not just completion.**

---

## Good Luck!

Clean data is the foundation of good models.

Take time to understand each preprocessing step.