

Introduction to Deep Learning with R

Intermediate Predictive Analytics

Prof.Asc. Endri Raco, Ph.D.

Department of Mathematical Engineering
Polytechnic University of Tirana

November 2025

Section 1

Section 1: Introduction

Why Deep Learning?

Deep learning enables machines to **automatically discover patterns** in data without extensive manual feature engineering.

Key Applications:

- Image recognition and computer vision
- Natural language processing
- Time series forecasting
- Customer behavior prediction
- Recommendation systems

Today's Focus: Understanding neural network fundamentals and implementing them in R using the keras package.

Course Context

This lecture fits into our course progression:

Previous Topics:

- Basetable construction
- Feature engineering
- Traditional ML models

Today:

- Neural networks
- Deep learning
- Keras in R

Next Topics:

- Advanced architectures
- Model optimization
- Deployment

Learning Objectives

By the end of this lecture, you will be able to:

- ① Understand neural network architecture and forward propagation
- ② Implement activation functions in R
- ③ Build multi-layer neural networks
- ④ Use keras to train and evaluate models
- ⑤ Recognize when deep learning adds value

Section 2

Section 2: Neural Network Fundamentals

The Core Problem: Feature Interactions

Scenario: Predict bank transactions based on customer characteristics

Table 1: Sample Customer Data

Customer	Children	Accounts	Balance	Retired
A	0	1	5000	No
B	2	3	50000	Yes
C	3	2	15000	No

The Challenge: How do retired customers with high balances differ from young families with multiple accounts?

Linear Regression vs Neural Networks

Linear Regression:

- Each feature contributes independently
- Interactions must be manually specified
- Limited expressiveness
- Formula:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots$$

Neural Networks:

- Automatically learns interactions
- Hidden layers discover patterns
- Highly flexible
- Can model complex nonlinear relationships

Key Insight: Neural networks create intermediate representations that capture feature interactions automatically.

Neural Network Architecture

Basic Components:

- ① **Input Layer:** Raw features (e.g., age, balance, accounts)
- ② **Hidden Layer(s):** Learned intermediate representations
- ③ **Output Layer:** Final predictions
- ④ **Weights:** Parameters learned during training
- ⑤ **Activation Functions:** Introduce nonlinearity

Key Terminology:

- **Neurons/Nodes:** Individual computational units
- **Layers:** Collections of neurons
- **Deep Learning:** Multiple hidden layers (depth)


Simple Network Example

Consider a network to predict transactions:

Architecture:

Input (4 features) → Hidden (2 nodes) → Output (1 value)

Simple Neural Network Architecture



Section 3

Section 3: Forward Propagation

What is Forward Propagation?

Forward propagation is the process of computing predictions by passing input data through the network layer by layer.

The Algorithm:

- ① Start with input values
- ② For each layer:
 - Multiply inputs by weights
 - Sum the weighted inputs
 - Apply activation function
 - Pass result to next layer
- ③ Output final prediction

This is how neural networks make predictions for a single observation.

Mathematical Notation

For a simple two-layer network:

Input to Hidden Layer:

$$h_j = f \left(\sum_{i=1}^n w_{ij}^{(1)} x_i + b_j^{(1)} \right)$$

where:

- x_i are input features
- $w_{ij}^{(1)}$ are weights from input i to hidden node j
- $b_j^{(1)}$ is the bias term
- f is the activation function
- h_j is the hidden layer activation

Hidden to Output Layer:

$$\hat{y} = g \left(\sum_{i=1}^m w_{ji}^{(2)} h_i + b_j^{(2)} \right)$$

Forward Propagation Example: Setup

Given: A customer with 2 children and 3 accounts

```
# Input data  
input_data <- c(children = 2, accounts = 3)  
  
# Weights for first hidden node  
weights_node_0 <- c(1, 1)  
  
# Weights for second hidden node  
weights_node_1 <- c(-1, 1)  
  
# Weights for output layer  
weights_output <- c(2, -1)
```

Task: Compute the predicted number of transactions

Forward Propagation: Hidden Layer

Computing Hidden Node 0:

```
# Weighted sum for node 0
node_0_input <- sum(input_data * weights_node_0)
cat("Node 0 input:", node_0_input, "\n")
```

Node 0 input: 5

```
# No activation yet (linear)
node_0_value <- node_0_input
```

Computing Hidden Node 1:

```
# Weighted sum for node 1
node_1_input <- sum(input_data * weights_node_1)
cat("Node 1 input:", node_1_input, "\n")
```

Node 1 input: 1

```
# No activation yet (linear)
```

Forward Propagation: Output Layer

Computing Final Prediction:

```
# Combine hidden layer outputs
hidden_layer <- c(node_0_value, node_1_value)
cat("Hidden layer values:", hidden_layer, "\n")

## Hidden layer values: 5 1

# Compute final output
prediction <- sum(hidden_layer * weights_output)
cat("\nFinal prediction:", prediction, "transactions\n")

## 
## Final prediction: 9 transactions
```

The network predicts **9 transactions** based on our manually specified weights.

Generalizing: Forward Propagation Function

```
# Display the forward propagation function
forward_propagate

## function (input_data, weights)
## {
##     node_0 <- sum(input_data * weights$node_0)
##     node_1 <- sum(input_data * weights$node_1)
##     hidden_layer <- c(node_0, node_1)
##     output <- sum(hidden_layer * weights$output)
##     return(output)
## }
```

Testing Forward Propagation Function

```
# Define weights structure
weights <- list(
  node_0 = c(1, 1),
  node_1 = c(-1, 1),
  output = c(2, -1)
)

# Test the function
result <- forward_propagate(input_data, weights)
cat("Prediction:", result, "\n")

## Prediction: 9
```

Section 4

Section 4: Activation Functions

The Need for Nonlinearity

Problem: Without activation functions, neural networks collapse to linear models.

Consider stacking two linear transformations:

$$y = W_2(W_1x) = (W_2W_1)x = W_{combined}x$$

This is equivalent to a single linear transformation!

Solution: Apply nonlinear activation functions after each layer:

$$y = f_2(W_2f_1(W_1x))$$

Now the network can learn complex, nonlinear patterns.

Common Activation Functions

1. ReLU (Rectified Linear Unit) - Most Popular

$$f(x) = \max(0, x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

2. Tanh (Hyperbolic Tangent)

$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

3. Sigmoid (Logistic)

$$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$

Implementing Activation Functions

```
# Display activation functions
```

```
cat("ReLU function:\n")
```

```
## ReLU function:
```

```
relu
```

```
## function (x)
```

```
## {
```

```
##     pmax(0, x)
```

```
## }
```

```
cat("\nSigmoid function:\n")
```

```
##
```

```
## Sigmoid function:
```

```
sigmoid
```

```
## function (x)
```


Testing Activation Functions

```
# Test the functions
test_values <- c(-2, -1, 0, 1, 2)

results <- data.frame(
  Input = test_values,
  ReLU = relu(test_values),
  Tanh = tanh(test_values),
  Sigmoid = sigmoid(test_values)
)
kable(results, digits = 3, booktabs = TRUE)
```

Input	ReLU	Tanh	Sigmoid
-2	0	-0.964	0.119
-1	0	-0.762	0.269
0	0	0.000	0.500
1	1	0.762	0.731

Visualizing Activation Functions



Activation Function Properties

Function	Range	Zero-Centered	Computational Cost	Typical Use
ReLU	[0, Inf)	No	Low	Hidden layers
Tanh	(-1, 1)	Yes	Medium	Hidden layers
Sigmoid	(0, 1)	No	Medium	Output (binary)

Rule of thumb: Use ReLU for hidden layers unless you have a specific reason not to.

Forward Propagation with ReLU

```
# Display the function with activation
forward_propActivated

## function (input_data, weights, activation = relu)
## {
##     node_0_input <- sum(input_data * weights$node_0)
##     node_0_output <- activation(node_0_input)
##     node_1_input <- sum(input_data * weights$node_1)
##     node_1_output <- activation(node_1_input)
##     hidden_layer <- c(node_0_output, node_1_output)
##     output <- sum(hidden_layer * weights$output)
##     return(list(hidden = hidden_layer, output = output))
## }
```

Comparing Linear vs ReLU Networks

```
# Without activation (linear)
linear_result <- forward_propagate(input_data, weights)

# With ReLU activation
relu_result <- forward_prop_activate(input_data, weights, relu)

cat("Linear network output:", linear_result, "\n")

## Linear network output: 9

cat("ReLU network output:", relu_result$output, "\n")

## ReLU network output: 9

cat("Hidden layer activations:", relu_result$hidden, "\n")

## Hidden layer activations: 5 1
```

Notice how ReLU sets negative values to zero, creating sparse

Section 5

Section 5: Deeper Networks

Why Add More Layers?

Single Hidden Layer Networks:

- Can theoretically approximate any function (Universal Approximation Theorem)
- May require exponentially many neurons
- Difficult to learn hierarchical patterns

Multiple Hidden Layers:

- Learn hierarchical representations
- More parameter efficient
- Better generalization on complex tasks
- Each layer captures different abstraction levels

Hierarchical Feature Learning

Image Recognition Example:

Layer	Learns
Layer 1	Edges, corners, basic textures
Layer 2	Simple shapes, object parts
Layer 3	Object components (eyes, wheels)
Layer 4	Complete objects (faces, cars)
Output	Object categories

Each layer builds upon representations from previous layers.

Multi-Layer Network Implementation

```
# Weights for a 2-hidden-layer network
multilayer_weights <- list(
  # Input (2) to Hidden Layer 1 (2 nodes)
  h1_node0 = c(2, 4),
  h1_node1 = c(-5, -1),
  # Hidden Layer 1 (2) to Hidden Layer 2 (2 nodes)
  h2_node0 = c(-1, 1),
  h2_node1 = c(2, 4),
  # Hidden Layer 2 (2) to Output (1)
  output = c(-3, 7)
)
```

Multi-Layer Forward Propagation

```
# Display the multi-layer function
multilayer_forward

## function (input_data, weights)
## {
##     h1_0 <- relu(sum(input_data * weights$h1_node0))
##     h1_1 <- relu(sum(input_data * weights$h1_node1))
##     hidden1 <- c(h1_0, h1_1)
##     h2_0 <- relu(sum(hidden1 * weights$h2_node0))
##     h2_1 <- relu(sum(hidden1 * weights$h2_node1))
##     hidden2 <- c(h2_0, h2_1)
##     output <- sum(hidden2 * weights$output)
##     return(list(h1 = hidden1, h2 = hidden2, output = output))
## }
```

Testing Multi-Layer Network

```
# Test the multi-layer network
result <- multilayer_forward(input_data, multilayer_weights)

cat("Hidden Layer 1:", result$h1, "\n")

## Hidden Layer 1: 16 0

cat("Hidden Layer 2:", result$h2, "\n")

## Hidden Layer 2: 0 32

cat("Output:", result$output, "\n")

## Output: 224
```

Network Depth Analysis

Table 4: Network Depth Comparison

Architecture	Output	Parameters
No Hidden Layers	5	2
1 Hidden Layer	9	6
2 Hidden Layers	224	10

Deeper networks can learn more complex representations with similar parameter counts.

Section 6

Section 6: Building Networks with Keras

Introduction to Keras

Keras is a high-level neural networks API that runs on TensorFlow.

Installation (one-time setup):

```
install.packages("keras")
library(keras)
install_keras() # Installs TensorFlow backend
```

Why Keras?

- User-friendly API
- Fast prototyping
- Production-ready
- Extensive documentation
- Large community support

Creating a Simple Keras Model

```
library(keras)

# Define model architecture
model <- keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu",
              input_shape = c(4)) %>%
  layer_dense(units = 32, activation = "relu") %>%
  layer_dense(units = 1)

# View model structure
summary(model)
```

Architecture:

- Input: 4 features
- Hidden layer 1: 64 neurons, ReLU
- Hidden layer 2: 32 neurons, ReLU

Generating Synthetic Bank Data

```
# Create synthetic customer dataset
set.seed(123)
n <- 1000


bank_data <- tibble(
  children = sample(0:4, n, replace = TRUE),
  accounts = sample(1:5, n, replace = TRUE),
  balance = rnorm(n, 25000, 15000),
  age = sample(25:75, n, replace = TRUE),
  retired = ifelse(age > 65, 1, 0)
) %>%
  mutate(
    # Simulate complex interaction for transactions
    transactions = 20 +
      5 * children +
      8 * accounts +
      0.0003 * balance +
```

Exploring the Data

Table 5: Bank Customer Data Summary

children	accounts	balance	age	retired	transactions
Min.	Min.	Min.	Min.	Min.	Min. :
:0.000	:1.000	:-20718	:25.00	:0.000	14.31
1st	1st	1st Qu.:	1st	1st	1st Qu.:
Qu.:1.000	Qu.:2.000	15215	Qu.:37.00	Qu.:0.000	48.24
Median	Median	Median :	Median	Median	Median :
:2.000	:3.000	24646	:50.00	:0.000	67.41
Mean	Mean	Mean :	Mean	Mean	Mean :
:1.985	:2.987	24931	:50.02	:0.206	71.05
3rd	3rd	3rd Qu.:	3rd	3rd	3rd Qu.:
Qu.:3.000	Qu.:4.000	35420	Qu.:63.00	Qu.:0.000	90.68
Max.	Max.	Max. :	Max.	Max.	Max.
:4.000	:5.000	74358	:75.00	:1.000	:153.37

Data Distribution



Data Preprocessing

```
# Separate features and target
X <- bank_data %>%
  select(children, accounts, balance, age, retired) %>%
  as.matrix()

y <- bank_data$transactions

# Standardize features (critical for neural networks!)
X_scaled <- scale(X)

# Train/test split (80/20)
set.seed(42)
train_idx <- sample(1:n, 0.8 * n)

X_train <- X_scaled[train_idx, ]
y_train <- y[train_idx]
```

Building the Keras Model

```
library(keras)

# Define architecture
model <- keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu",
              input_shape = c(ncol(X_train))) %>%
  layer_dense(units = 32, activation = "relu") %>%
  layer_dense(units = 1)

# Compile model
model %>% compile(
  optimizer = "adam",                      # Adaptive learning rate
  loss = "mse",                            # Mean squared error
  metrics = c("mae")                        # Mean absolute error
)

# View architecture
```

Model Architecture Summary

When you run `summary(model)`, you'll see:

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	384
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33
Total params		2,497

Parameter calculation:

$$(5 \times 64 + 64) + (64 \times 32 + 32) + (32 \times 1 + 1) = 2,497$$

Training the Model

```
# Train the model
history <- model %>% fit(
  x = X_train,
  y = y_train,
  epochs = 50,                      # Number of complete passes
  batch_size = 32,                   # Samples per gradient update
  validation_split = 0.2,            # Use 20% for validation
  verbose = 1                        # Show progress
)

# Plot training history
plot(history)
```

Training Parameters Explained

Key hyperparameters:

- **epochs**: Number of complete passes through training data
 - Too few: underfitting
 - Too many: overfitting
- **batch_size**: Samples processed before weight update
 - Smaller: more updates, noisier gradients
 - Larger: fewer updates, smoother gradients
- **validation_split**: Fraction for monitoring overfitting
 - Typical: 0.1 to 0.3

Making Predictions

```
# Generate predictions on test set
predictions <- model %>% predict(X_test)

# Evaluate model
evaluation <- model %>% evaluate(X_test, y_test, verbose = 0)

cat("Test MSE:", evaluation$loss, "\n")
cat("Test MAE:", evaluation$mae, "\n")

# Create results dataframe
results <- tibble(
  Actual = y_test,
  Predicted = as.vector(predictions)
)

head(results, 10)
```

Visualizing Model Performance

```
# Predicted vs Actual plot
ggplot(results, aes(x = Actual, y = Predicted)) +
  geom_point(alpha = 0.5, color = "steelblue") +
  geom_abline(intercept = 0, slope = 1,
              color = "red", linetype = "dashed", linewidth =
  labs(title = "Model Performance: Predicted vs Actual",
       x = "Actual Transactions",
       y = "Predicted Transactions") +
  coord_fixed() +
  theme_minimal()
```

Residual Analysis

```
# Calculate residuals
results <- results %>%
  mutate(Residual = Actual - Predicted)

# Residual plot
ggplot(results, aes(x = Predicted, y = Residual)) +
  geom_point(alpha = 0.5) +
  geom_hline(yintercept = 0, color = "red",
             linetype = "dashed") +
  labs(title = "Residual Plot",
       x = "Predicted Values",
       y = "Residuals (Actual - Predicted)") +
  theme_minimal()
```

Section 7

Section 7: Best Practices

When to Use Deep Learning

Deep learning excels when:

- Large datasets ($n > 10,000$)
- Complex nonlinear relationships
- High-dimensional inputs
- Feature engineering is difficult
- Images, text, or sequential data

Consider alternatives when:

- Small datasets ($n < 1,000$)
- Simple linear relationships
- Interpretability is critical
- Limited computational resources
- Fast training is essential

Neural Network Best Practices

1. Always normalize input features

- Use `scale()` in R or keras preprocessing layers
- Neural networks are sensitive to feature scales

2. Start simple, then add complexity

- Begin with 1-2 hidden layers
- Add layers/neurons only if needed

3. Use appropriate activation functions

- Hidden layers: ReLU (default)
- Binary output: sigmoid
- Multi-class output: softmax
- Regression output: none (linear)

Training Best Practices

4. Monitor training and validation loss

- Both should decrease together
- Divergence indicates overfitting
- Use early stopping

5. Use appropriate batch sizes

- Too small: unstable training
- Too large: poor generalization
- Typical: 32, 64, 128

6. Tune learning rate carefully

- Too high: divergence
- Too low: slow convergence
- Use adaptive optimizers (Adam)

Common Pitfalls to Avoid

- ① **Forgetting to normalize features** - Leads to poor convergence
- ② **Using too many parameters** - Causes overfitting on small datasets
- ③ **Not using validation data** - Cannot detect overfitting
- ④ **Training for too many epochs** - Overfits to training noise
- ⑤ **Inappropriate activation for output** - Wrong predictions
- ⑥ **Not setting random seeds** - Non-reproducible results

Section 8

Section 8: Summary and Practice

Key Concepts Review

Core Concepts:

- ① Neural networks learn feature interactions automatically
- ② Forward propagation computes predictions layer-by-layer
- ③ Activation functions introduce essential nonlinearity
- ④ Deeper networks learn hierarchical representations
- ⑤ Keras provides a high-level API for building models

Practical Skills:

- ① Implementing forward propagation manually
- ② Using activation functions correctly
- ③ Building and training keras models
- ④ Evaluating model performance

Classwork Assignment

Task: Build a customer churn prediction model

```
# Customer data
customer <- c(
  age = 45,
  income = 75000,
  accounts = 3,
  children = 2,
  balance = 50000
)
```

Requirements:

- ① Implement manual forward propagation (2 hidden layers, ReLU)
- ② Output layer uses sigmoid for churn probability
- ③ Build equivalent keras model
- ④ Compare manual vs keras predictions

Assignment Details

Architecture:

- Input: 5 features (age, income, accounts, children, balance)
- Hidden Layer 1: 4 nodes, ReLU activation
- Hidden Layer 2: 3 nodes, ReLU activation
- Output: 1 node, sigmoid activation (churn probability)

Deliverables:

- ① R script with manual implementation
- ② Keras model code
- ③ Comparison of results
- ④ Brief interpretation (2-3 sentences)

Due: Next class session

Assignment Grading Rubric

Component	Points
Manual forward propagation (correct)	30
Proper activation function usage	20
Keras model architecture (correct)	20
Model training and evaluation	15
Code quality and documentation	10
Interpretation of results	5
Total	100

Additional Resources

Essential Reading:

- *Deep Learning with R* (Chollet & Allaire) - Chapters 1-4
- *Neural Networks and Deep Learning* (Nielsen) - Free online
- Keras R documentation: keras.rstudio.com

Online Resources:

- TensorFlow tutorials: tensorflow.rstudio.com
- Fast.ai Practical Deep Learning course
- DeepLearning.AI specialization (Coursera)

Practice Datasets:

- MNIST digits (keras built-in)
- UCI Machine Learning Repository
- Kaggle competitions