

# Introduction to Deep Learning with R

## Intermediate Predictive Analytics

Prof.Asc. Endri Raco, Ph.D.

Department of Mathematical Engineering  
Polytechnic University of Tirana

November 2025

# Section 1

## Introduction

# Why Deep Learning?

Deep learning has revolutionized predictive analytics by enabling machines to automatically discover patterns in data without extensive feature engineering.

## Key Applications:

- Image recognition and computer vision
- Natural language processing
- Time series forecasting
- Customer behavior prediction
- Recommendation systems

**Today's Focus:** Understanding neural networks and implementing them in R using the keras package.

# Course Context

This lecture is part of our **Intermediate Predictive Analytics** course sequence:

**Previous Topics:** Basetable construction, feature engineering, traditional ML models

**Today:** Introduction to neural networks and deep learning

**Next Topics:** Advanced architectures, model optimization, deployment

## Learning Objectives:

- Understand neural network architecture and forward propagation
- Implement activation functions in R
- Build multi-layer networks
- Recognize when deep learning adds value

## Section 2

### Neural Networks: Core Concepts

# The Motivation: Capturing Interactions

Consider predicting bank transactions based on customer characteristics.

## Linear Regression View:

- Each feature contributes independently
- Interactions must be manually specified
- Limited ability to capture complex patterns

## Neural Network View:

- Automatically learns feature interactions
- Hidden layers create intermediate representations
- Can model highly nonlinear relationships

# Example: Bank Transaction Prediction

**Scenario:** Predict number of transactions next year

## Input Features:

- Number of children
- Number of existing accounts
- Bank balance
- Retirement status

**Question:** How do these features interact?

- Retired customers with high balances behave differently
- Multiple accounts interact with number of children
- Neural networks learn these patterns automatically

# Network Architecture: Building Blocks

## Components:

- **Input Layer:** Raw features
- **Hidden Layers:** Learned representations
- **Output Layer:** Predictions
- **Weights:** Learned parameters

## Key Insight:

- Each connection has a weight
- Weights are learned from data
- More layers = deeper networks
- “Deep” learning comes from multiple hidden layers

## Section 3

### Forward Propagation

# What is Forward Propagation?

**Forward propagation** is the process of computing predictions by passing input data through the network layer by layer.

## The Process:

- ① Start with input values
- ② Multiply by weights and sum (dot product)
- ③ Apply activation function
- ④ Pass result to next layer
- ⑤ Repeat until reaching output

**Critical:** This is how neural networks make predictions for a single data point.

## Simple Example: Two Inputs, One Hidden Layer

Consider inputs: number of children = 2, number of accounts = 3

```
# Input data
input_data <- c(2, 3)

# Weights for hidden layer nodes
# First hidden node
weights_node_0 <- c(1, 1)
# Second hidden node
weights_node_1 <- c(-1, 1)

# Weights for output layer
weights_output <- c(2, -1)
```

# Computing Hidden Layer Values I

## Node 0 calculation:

```
# Compute weighted sum for node 0
node_0_value <- sum(input_data * weights_node_0)
node_0_value
```

```
## [1] 5
```

# Computing Hidden Layer Values II

## Node 1 calculation:

```
# Compute weighted sum for node 1
node_1_value <- sum(input_data * weights_node_1)
node_1_value

## [1] 1
```

The hidden layer now has values: 5 and 1

# Computing Final Output

```
# Combine hidden layer values
hidden_layer_values <- c(node_0_value, node_1_value)

# Compute final output
output <- sum(hidden_layer_values * weights_output)
output

## [1] 9
```

**Result:** The network predicts 9 transactions

This is pure forward propagation without activation functions (coming next!)

# Complete Forward Propagation Function I

```
# Function to perform forward propagation
forward_prop <- function(input_data, weights) {
  # Calculate hidden layer
  node_0_value <- sum(input_data * weights$node_0)
  node_1_value <- sum(input_data * weights$node_1)

  # Combine hidden layer
  hidden_layer <- c(node_0_value, node_1_value)

  # Calculate output
  output <- sum(hidden_layer * weights$output)

  return(output)
}
```

## Complete Forward Propagation Function II

```
# Define weights structure
weights <- list(
  node_0 = c(1, 1),
  node_1 = c(-1, 1),
  output = c(2, -1)
)

# Test the function
result <- forward_prop(input_data, weights)
cat("Prediction:", result, "\n")

## Prediction: 9
```

## Section 4

### Activation Functions

# Why Activation Functions?

**Problem:** Without activation functions, neural networks are just linear models

- Stacking linear transformations = one big linear transformation
- Cannot capture nonlinear patterns
- Limited expressiveness

**Solution:** Apply nonlinear activation functions

- Introduces nonlinearity into the network
- Enables learning of complex patterns
- Each layer can learn increasingly abstract features

# Common Activation Functions I

## 1. ReLU (Rectified Linear Unit)

- Most popular in deep learning
- $f(x) = \max(0, x)$
- Fast to compute
- Helps avoid vanishing gradients

## 2. Tanh (Hyperbolic Tangent)

- $f(x) = \tanh(x)$
- Outputs in range  $[-1, 1]$
- Centered around zero

# Common Activation Functions II

## 3. Sigmoid

- $f(x) = \frac{1}{1+e^{-x}}$
- Outputs in range  $[0, 1]$
- Good for binary classification

## When to use what?

- **Hidden layers:** ReLU (default)
- **Output layer:** depends on task
  - Binary: sigmoid
  - Multi-class: softmax

# ReLU: The Default Choice I

**ReLU (Rectified Linear Unit)** is the most widely used activation function.

## ReLU: The Default Choice II

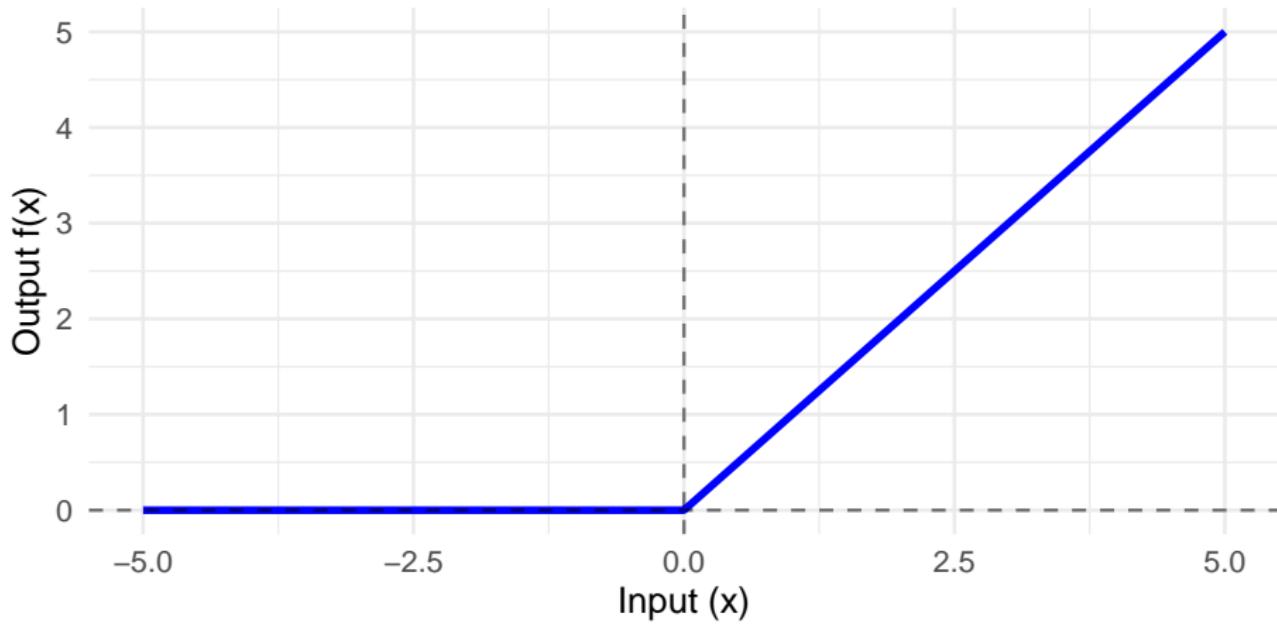
```
# Create data for visualization
x <- seq(-5, 5, by = 0.1)
y_relu <- pmax(0, x) # ReLU: max(0, x)

# Plot ReLU function
ggplot(data.frame(x = x, y = y_relu), aes(x, y)) +
  geom_line(color = "blue", linewidth = 1.2) +
  geom_hline(yintercept = 0, linetype = "dashed",
             alpha = 0.5) +
  geom_vline(xintercept = 0, linetype = "dashed",
             alpha = 0.5) +
  labs(title = "ReLU Activation Function",
       subtitle = "f(x) = max(0, x)",
       x = "Input (x)", y = "Output f(x)") +
  theme_minimal(base_size = 12)
```

# ReLU: The Default Choice III

ReLU Activation Function

$$f(x) = \max(0, x)$$



# Implementing Activation Functions in R I

```
# ReLU implementation
relu <- function(x) {
  # Element-wise maximum with 0
  pmax(0, x)
}
```

```
# Tanh is built-in: tanh(x)
```

```
# Sigmoid implementation
sigmoid <- function(x) {
  1 / (1 + exp(-x))
}
```

# Implementing Activation Functions in R II

```
# Test the functions
test_values <- c(-2, -1, 0, 1, 2)

results <- data.frame(
  Input = test_values,
  ReLU = relu(test_values),
  Tanh = tanh(test_values),
  Sigmoid = sigmoid(test_values)
)
print(results)
```

# Implementing Activation Functions in R III

```
##      Input ReLU      Tanh Sigmoid
## 1     -2    0 -0.9640276 0.1192029
## 2     -1    0 -0.7615942 0.2689414
## 3      0    0  0.0000000 0.5000000
## 4      1    1  0.7615942 0.7310586
## 5      2    2  0.9640276 0.8807971
```

# Forward Propagation with Activation Functions

```
# Enhanced forward propagation with ReLU
forward_prop_relu <- function(input_data, weights) {
  # Calculate hidden layer with ReLU activation
  node_0_input <- sum(input_data * weights$node_0)
  node_0_output <- relu(node_0_input)

  node_1_input <- sum(input_data * weights$node_1)
  node_1_output <- relu(node_1_input)

  # Combine hidden layer
  hidden_layer <- c(node_0_output, node_1_output)

  # Calculate output (no activation for regression)
  output <- sum(hidden_layer * weights$output)

  return(output)
}
```

# Comparing Activation Functions I

# Comparing Activation Functions II

```
# Compare multiple activation functions
x <- seq(-5, 5, by = 0.1)

comparison_data <- data.frame(
  x = x,
  ReLU = relu(x),
  Tanh = tanh(x),
  Sigmoid = sigmoid(x)
) %>%
  pivot_longer(cols = -x, names_to = "Function",
               values_to = "y")

ggplot(comparison_data, aes(x, y, color = Function)) +
  geom_line(linewidth = 1.2) +
  labs(title = "Comparison of Activation Functions",
       x = "Input (x)", y = "Output f(x)") +
  theme_minimal(base_size = 12) +
```

## Section 5

### Deeper Networks

# Why Go Deeper?

## Single hidden layer networks are limited

- Can approximate many functions (universal approximation theorem)
- But may require exponentially many neurons
- Difficult to learn hierarchical features

## Multiple hidden layers enable:

- ① **Hierarchical feature learning** - Each layer learns progressively abstract features
- ② **Improved generalization** - Better at capturing complex patterns
- ③ **Parameter efficiency** - Fewer total parameters for same expressiveness

# Example: Multi-Layer Network Architecture I

Consider a network with two hidden layers:

Input Layer → Hidden Layer 1 → Hidden Layer 2 → Output Layer

## Example: Multi-Layer Network Architecture II

```
# Define weights for two-hidden-layer network
multilayer_weights <- list(
  # Input to first hidden layer (2 nodes)
  hidden1_node0 = c(2, 4),
  hidden1_node1 = c(-5, -1),
  
  # First hidden to second hidden layer (2 nodes)
  hidden2_node0 = c(-1, 1),
  hidden2_node1 = c(2, 4),
  
  # Second hidden to output
  output = c(-3, 7)
)
```

# Forward Propagation Through Multiple Layers I

# Forward Propagation Through Multiple Layers II

```
multilayer_forward <- function(input_data, weights) {  
  # First hidden layer  
  h1_node0 <- relu(sum(input_data *  
                        weights$hidden1_node0))  
  h1_node1 <- relu(sum(input_data *  
                        weights$hidden1_node1))  
  hidden1_output <- c(h1_node0, h1_node1)  
  
  # Second hidden layer  
  h2_node0 <- relu(sum(hidden1_output *  
                        weights$hidden2_node0))  
  h2_node1 <- relu(sum(hidden1_output *  
                        weights$hidden2_node1))  
  hidden2_output <- c(h2_node0, h2_node1)  
  
  # Output layer  
  output <- sum(hidden2_output * weights$output)}
```

# Examining Network Activations

## Key Observations:

- ReLU sets negative values to zero
- Only some nodes activate for each input
- Creates sparse representations
- Different inputs activate different pathways

This **selective activation** is what enables the network to learn complex decision boundaries and feature hierarchies.

# Representation Learning

## What is Representation Learning?

Deep networks automatically discover useful representations of data in their hidden layers.

### Layer-by-layer learning:

- **Layer 1:** Simple patterns (edges, basic features)
- **Layer 2:** Combinations of simple patterns
- **Layer 3:** High-level concepts
- **Output:** Final prediction

### Benefits:

- Reduces need for manual feature engineering
- Learns task-specific representations
- Transfers knowledge between tasks
- Discovers patterns humans might miss

## Section 6

Building Neural Networks in R with Keras

# Introduction to Keras

Keras is a high-level neural networks API that runs on top of TensorFlow.

```
# Install keras (one-time setup)
install.packages("keras")
library(keras)
install_keras() # Installs TensorFlow backend

# Load keras for use
library(keras)
```

## Why Keras?

- Simple, consistent interface
- Quick prototyping
- Production-ready
- Extensive pre-trained models
- Active community

# Building Your First Neural Network I

```
# Define network architecture
model <- keras_model_sequential() %>%
  # First hidden layer: 100 units, ReLU activation
  layer_dense(units = 100, activation = 'relu',
              input_shape = c(10)) %>%
  # Second hidden layer: 100 units, ReLU activation
  layer_dense(units = 100, activation = 'relu') %>%
  # Output layer: 1 unit (regression)
  layer_dense(units = 1)

# View model architecture
summary(model)
```

# Building Your First Neural Network II

## Key Parameters:

- `units`: Number of neurons in the layer
- `activation`: Activation function to use
- `input_shape`: Dimensions of input (first layer only)

# Complete Example: Bank Transaction Prediction I

## Complete Example: Bank Transaction Prediction II

```
# Load and prepare data
library(tidyverse)

# Simulated bank customer data
set.seed(123)
n_customers <- 1000

customer_data <- tibble(
  n_children = sample(0:4, n_customers,
                      replace = TRUE),
  n_accounts = sample(1:5, n_customers,
                      replace = TRUE),
  balance = rnorm(n_customers, 10000, 5000),
  retired = sample(0:1, n_customers,
                   replace = TRUE),
  age = sample(25:75, n_customers, replace = TRUE)
) %>%
```

# Preprocessing and Train/Test Split I

# Preprocessing and Train/Test Split II

```
# Separate features and target
X <- customer_data %>%
  select(-n_transactions) %>%
  as.matrix()

y <- customer_data$n_transactions

# Normalize features (important for neural networks!)
X_normalized <- scale(X)

# Train/test split (80/20)
set.seed(42)
train_indices <- sample(1:n_customers,
                       0.8 * n_customers)

X_train <- X_normalized[train_indices, ]
y_train <- y[train_indices]
```

# Building and Compiling the Model I

## Building and Compiling the Model II

```
# Build the model
model <- keras_model_sequential() %>%
  layer_dense(units = 64, activation = 'relu',
              input_shape = c(ncol(X_train))) %>%
  layer_dense(units = 32, activation = 'relu') %>%
  layer_dense(units = 1)

# Compile the model
model %>% compile(
  optimizer = 'adam',           # Optimization algorithm
  loss = 'mse',                 # Mean squared error
  metrics = c('mae'))          # Mean absolute error
)

# View architecture
summary(model)
```

# Training the Model

```
# Train the model
history <- model %>% fit(
  x = X_train,
  y = y_train,
  epochs = 50,                      # Number of passes through data
  batch_size = 32,                   # Samples per gradient update
  validation_split = 0.2,            # Use 20% for validation
  verbose = 1                        # Show progress
)
# Plot training history
plot(history)
```

## Key Parameters:

- **epochs**: How many times to iterate through entire dataset
- **batch\_size**: Number of samples to use per gradient update
- **validation\_split**: Fraction of training data to use for validation

# Making Predictions and Evaluating I

```
# Make predictions on test set
predictions <- model %>% predict(X_test)

# Evaluate model performance
evaluation <- model %>% evaluate(X_test, y_test)

cat("Test MSE:", evaluation$loss, "\n")
cat("Test MAE:", evaluation$mae, "\n")
```

# Making Predictions and Evaluating II

```
# Compare predictions to actual values
results <- tibble(
  Actual = y_test,
  Predicted = predictions[,1]
)

# Visualization
ggplot(results, aes(Actual, Predicted)) +
  geom_point(alpha = 0.5) +
  geom_abline(intercept = 0, slope = 1,
              color = "red", linetype = "dashed") +
  labs(title = "Predicted vs Actual Transactions",
       x = "Actual Transactions",
       y = "Predicted Transactions") +
  theme_minimal()
```

## Section 7

### Key Concepts Summary

# What We've Learned Today

## Core Concepts:

- ① Neural networks learn feature interactions automatically
- ② Forward propagation: computing predictions layer by layer
- ③ Activation functions introduce nonlinearity
- ④ Deeper networks learn hierarchical representations

## Practical Skills:

- ① Implementing forward propagation in R
- ② Using common activation functions
- ③ Building networks with keras
- ④ Training and evaluating models

Next class: Optimization, backpropagation, and hyperparameter tuning

# When to Use Deep Learning

## Deep learning excels when:

- Large amounts of data available
- Complex nonlinear patterns exist
- Feature engineering is difficult
- Working with images, text, or sequences
- High-dimensional input data

## Consider alternatives when:

- Small datasets ( $n < 1000$ )
- Simple linear relationships
- Interpretability is critical
- Limited computational resources
- Fast training is essential

**Remember:** Deep learning is a powerful tool, but not always the best tool.

# Best Practices for Neural Networks

## ① Always normalize/standardize your input features

- Neural networks are sensitive to feature scales
- Use `scale()` in R or preprocessing layers in keras

## ② Start with simple architectures

- Begin with 1-2 hidden layers
- Add depth only if performance plateaus

## ③ Use ReLU activation in hidden layers

- Default choice for most problems
- Helps avoid vanishing gradient problem

## ④ Monitor training and validation loss

- Early stopping prevents overfitting
- Validation loss should decrease with training loss

## ⑤ Use appropriate output activation

- Regression: no activation (or linear)
- Binary classification: sigmoid
- Multi-class: softmax

## Section 8

Practice and Next Steps

# Classwork Assignment

**Task:** Implement forward propagation for a custom network

```
# Given data
customer <- c(
  age = 45,
  income = 75000,
  accounts = 3,
  children = 2,
  balance = 50000
)
```

## Network architecture:

- Input: 5 features
- Hidden layer 1: 4 nodes (ReLU)
- Hidden layer 2: 3 nodes (ReLU)
- Output: 1 node (churn probability with sigmoid)

# Additional Resources

## Essential Reading:

- Deep Learning with R (Chollet & Allaire) - Chapters 1-3
- Neural Networks and Deep Learning (Nielsen) - Free online book
- Keras R documentation: <https://keras.rstudio.com>

## Online Tutorials:

- RStudio keras tutorials and examples
- Fast.ai Practical Deep Learning course
- DeepLearning.AI courses on Coursera

## Practice Datasets:

- MNIST handwritten digits (keras built-in)
- UCI Machine Learning Repository
- Kaggle competition datasets